

CH08-320201

Algorithms and Data Structures

Lecture 7/8 — 27 Feb 2018

Prof. Dr. Michael Sedlmair

Jacobs University
Spring 2018

This & that

- “homework is too hard”
 - Strategies:
 - background reading (the Cormen book gives background to most of the things that we are talking about in class)
 - go to tutorials and ask questions
 - use the Moodle discussion forum
 - What will change:
 - Me: will try to have a closer connection between level of difficulty/rigor between class and homework
 - A little bit of independent thinking will remain to be necessary for assignments (important for learning effect)
 - in general —> tell me/us your feedback, it's not supposed to be a “one man show”
- medical excuse policy refined:
 - <http://vis.jacobs-university.de/teaching/ads/18s/>

2. Sorting and Searching

Recall: Sorting Problem

- Input:
 - sequence $\langle a_1, a_2, \dots, a_n \rangle$ of numbers.
- Output:
 - permutation $\langle a'_1, a'_2, \dots, a'_n \rangle$
 - such that $a'_1 \leq a'_2 \leq \dots \leq a'_n$

Recall: Insertion & Merge Sort

Time complexity:

	Insertion Sort	Merge Sort
Best Case	$\Theta(n)$	$\Theta(n \lg n)$
Average Case	$\Theta(n^2)$	$\Theta(n \lg n)$
Worst Case	$\Theta(n^2)$	$\Theta(n \lg n)$

Visualizations:

<http://www.sorting-algorithms.com/insertion-sort>

<http://www.sorting-algorithms.com/merge-sort>

How about storage space complexity?

In situ

- Definition:
 - In-situ algorithms refer to algorithms that operate with $\Theta(1)$ memory.
- In-situ sorting:
 - Sorting algorithms that need only a constant number of additional storings.
- Insertion Sort:
 - In-situ sorting.
- Merge Sort:
 - Not in-situ sorting.

2.1 Heap Sort

Motivation


- Try to develop an in-situ sorting algorithm with asymptotic runtime $\Theta(n \lg n)$.
- Use a sophisticated data structure to support the computations.

Data structure: Heap

Defintion.


A (binary) heap data-structure is an array which can be viewed as a nearly complete binary tree:

each level is completely full except possibly the last level, which is filled from left to right.



Heap as array

A heap can be stored as an array:




Heap as array

The array A representing the heap has two attributes:

- $A.length$
- $A.heapsize$

such that $0 \leq A.heapsize \leq A.length$.

There are only $A.heapsize$ valid elements of the heap.

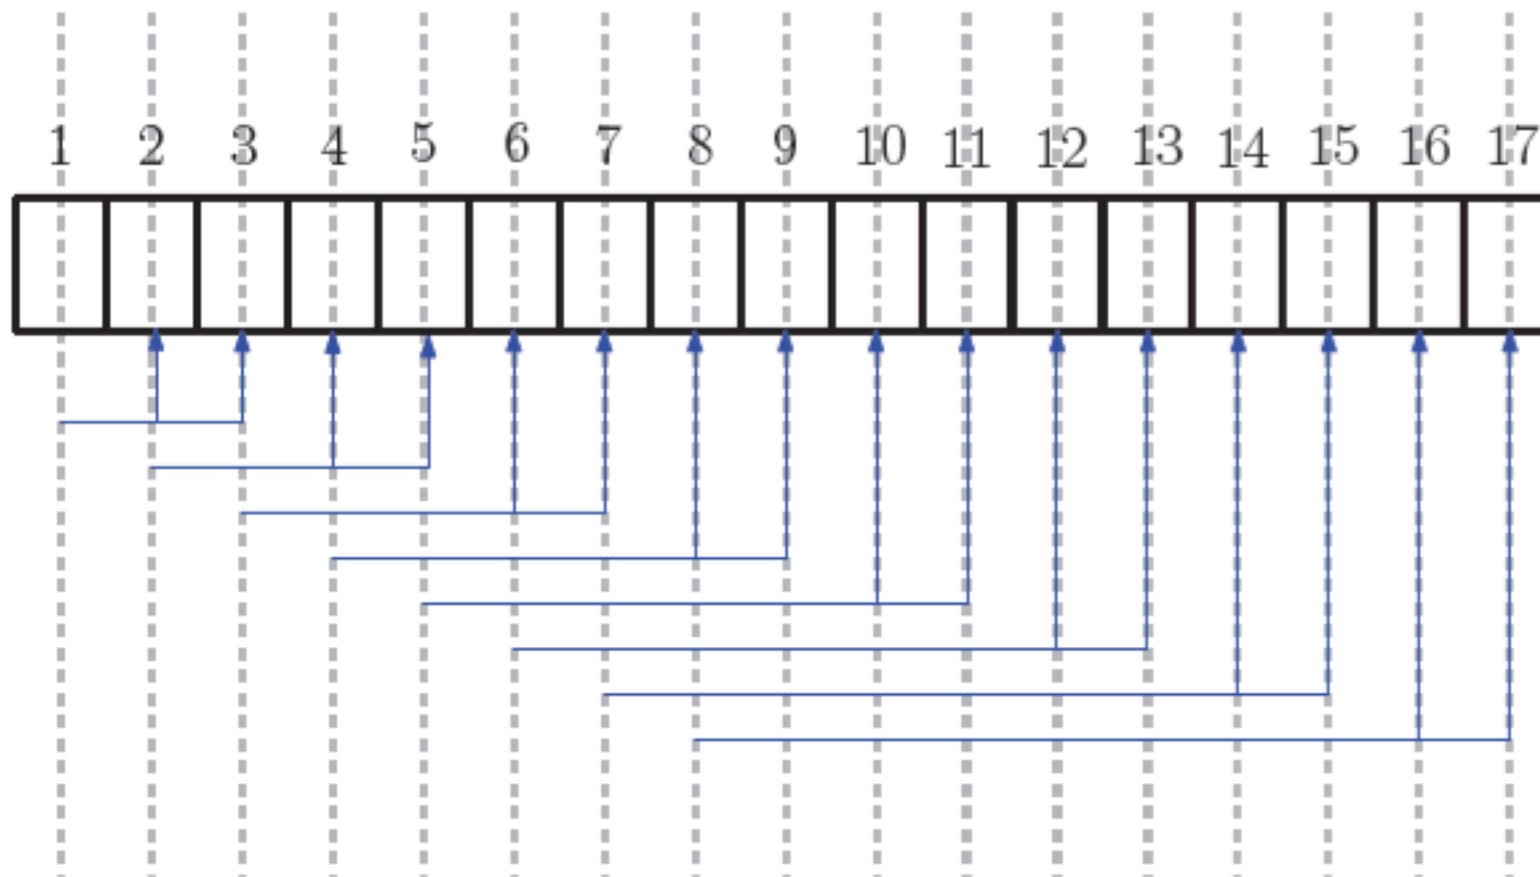


$A[1]$ is the root of the heap (root of the binary tree).

Heap as array

Given index i of an element of A , we can calculate:



- Parent(i): return $\text{floor}(i/2)$; // Right shift by 1 bit.
- Left(i): return $2i$; // Left shift by 1 bit.
- Right(i): return $2i + 1$; // Left shift by 1 bit
// and set LSB to 1



Max-heap property


In a max-heap, for every node i (other than the root),

$$A[\text{Parent}(i)] \geq A[i].$$



Recall: Height of a rooted tree

- The height of a node x is the length of the longest simple downward path from x to a leaf.
- The height of a rooted tree is the height of its root.



Heap height

Theorem:

A heap with n elements has height $h = \lfloor \lg n \rfloor$.

Heap height

Proof:

Heap height h implies that there are $h + 1$ levels (levels 0 to h).

As a heap is a nearly complete binary tree,
the last guaranteed complete level is level $h - 1$.

The level h may be incomplete but has at least one element.

Number of elements in complete levels 0 to $h - 1$ is

$$1 + 2 + 2^2 + \dots + 2^{h-1} = 2^h - 1.$$

So, $n > 2^h - 1$ or (since it is an integer) $n \geq 2^h$.

If all levels 0 to h were complete, the number of elements would be $2^{h+1} - 1$. So, $n \leq 2^{h+1} - 1$.

Heap height

Proof (continued):

Combining the two inequalities:

$$2^h \leq n \leq 2^{h+1} - 1$$

As $2^{h+1} > 2^{h+1} - 1 \geq 2^h$ for $h \geq 0$,

$$h+1 > \lg(2^{h+1} - 1) \geq h$$

Thus, $\lg(2^{h+1} - 1) = h + a$ with $a \in [0, 1)$,


which leads to $h \leq \lg n \leq h + a$ with $a \in [0, 1)$.

Hence, $h = \lfloor \lg n \rfloor$.


Max-Heapify (A, i)

Precondition:


When Max-Heapify (A, i) is called,
binary-trees rooted at Left(i) and Right(i) are valid
max-heaps, but $A[i]$ may be smaller than its children.




Max-Heapify (A,i)



Max-Heapify (A,i)



Max-Heapify (A,i)



Max-Heapify (A, i)

MAX-HEAPIFY(A, i)

```
1   $l = \text{LEFT}(i)$ 
2   $r = \text{RIGHT}(i)$ 
3  if  $l \leq A.\text{heap-size}$  and  $A[l] > A[i]$ 
4       $largest = l$ 
5  else  $largest = i$ 
6  if  $r \leq A.\text{heap-size}$  and  $A[r] > A[largest]$ 
7       $largest = r$ 
8  if  $largest \neq i$ 
9      exchange  $A[i]$  with  $A[largest]$ 
10     MAX-HEAPIFY( $A, largest$ )
```

Max-Heapify (A,i)

Complexity:

It is $O(h) = O(\lg n)$,

as in the worst case the new element has to go down all the way to the last level.

Converting an existing array A to a max-heap

- Call MAX-HEAPIFY, on what set?
 - all inner nodes
 - > float them down if necessary (MAX-HEAPIFY)
 - not necessary to call on leaves node

Leaves of a heap of size n

Let $n = A.\text{heapsize}$.

Where is the parent of the last element of a heap?

At index $\lfloor n/2 \rfloor$.

Therefore, the element at index $\lfloor n/2 \rfloor + 1$

does not have a child in the heap, and hence is a leaf.

In a heap, there are $\lfloor n/2 \rfloor$ leaves:

from index $\lfloor n/2 \rfloor + 1$ to n .

Each leaf is the root of a valid max-heap of size 1.

Converting an existing array A to a max-heap

BUILD-MAX-HEAP(A)

```
1   $A.heap-size = A.length$   
2  for  $i = \lfloor A.length/2 \rfloor$  downto 1  
3      MAX-HEAPIFY( $A, i$ )
```

Loop invariant:

At the start of each iteration of the for loop,
each node $i + 1, \dots, n$ is the root of a max-heap.

Converting an existing array A to a max-heap

Initialization: $i = \lfloor n/2 \rfloor$.

Nodes with indices $> i$ are leaves.

Maintenance:

At iteration i , nodes with indices $i + 1, \dots, n$ are roots of valid max-heaps, but the i -th node is not (necessarily).

But this is exactly the precondition of Max-Heapify.


Termination: $i = 0$.

Node $i = 1$ is the root of a valid max-heap.


Build-max-heap

A


4	1	3	2	16	9	10	14	8	7
---	---	---	---	----	---	----	----	---	---




Build-max-heap




Build-max-heap




Build-max-heap



Build-max-heap



Build-max-heap



Build-max-heap

What is the time complexity?

Theorem:

Let m_h be the number of nodes of height h in any n element heap $T(n)$.

$$\text{Then, } m_h(T, n) \leq \left\lceil \frac{n}{2^{h+1}} \right\rceil$$

(Proof by induction over h)

Build-max-heap

Time complexity:

Time needed by Max-Heapify when called on a node of height h is $O(h)$.

Therefore, total cost of Build-Max-Heap(A) is upper bounded by

$$\begin{aligned}\sum_{h=0}^{\lfloor \lg n \rfloor} \left\lceil \frac{n}{2^{h+1}} \right\rceil O(h) &= O\left(n \sum_{h=0}^{\lfloor \lg n \rfloor} \frac{h}{2^h}\right) \\ &= O\left(n \sum_{h=0}^{\infty} \frac{h}{2^h}\right) \\ &= O(n).\end{aligned}$$

$$\text{using } \sum_{k=0}^{\infty} k x^k = \frac{x}{(1-x)^2} \text{ if } |x| < 1.$$

Build-max-heap

Time complexity:

Conclusion:

We can convert an unordered array into a max-heap in linear time.

Heap Sort


- Start by generating a max-heap.
- The maximum element of a max-heap is at the root.
- Put it in its right sorted place at $n = A.\text{heapsize}$ by swapping it with the last element $A[n]$, which now becomes $A[1]$.
- Decrement the heap size to create a smaller heap and thus implicitly remove the last element (the maximum) from the heap.
- The new $A[1]$ may not satisfy the max-heap property, so float it down.
- Iterate.

Heapsort


HEAPSORT(A)

```
1  BUILD-MAX-HEAP( $A$ )
2  for  $i = A.length$  downto 2
3      exchange  $A[1]$  with  $A[i]$ 
4       $A.heap-size = A.heap-size - 1$ 
5      MAX-HEAPIFY( $A, 1$ )
```


Heapsort




Heapsort




Heapsort




Heapsort




Heapsort




Heapsort




Heapsort




Heapsort



Heapsort



Heapsort



Runtime Analysis

HEAPSORT(*A*)

```
1  BUILD-MAX-HEAP(A)
2  for i = A.length downto 2
3      exchange A[1] with A[i]
4      A.heap-size = A.heap-size − 1
5      MAX-HEAPIFY(A, 1)
```

- RuntimeCosts:
 $O(n) + O(n \lg n) = O(n \lg n)$
- MemoryCosts:
 $O(1)$, i.e., in-situ sorting
- Visualization:
<http://www.sorting-algorithms.com/heap-sort>

Heap as a data structure

- Heaps are a data structure that can be used for other purposes, as well.
- In particular, a max-heap is often used to build a max-priority queue.

Max-priority queues

Definition:

A priority queue is a data structure for maintaining a set S of elements, each with an associated value called a key.

Definition:

A max-priority queue is a priority queue that supports the following operations:

- $\text{Maximum}(S)$: return element from S with largest key.
- $\text{Extract-Max}(S)$: remove and return element from S with largest key.
- $\text{Increase-Key}(S, x, k)$: increase the value of the key of element x to k , where k is assumed to be larger or equal than the current key.
- $\text{Insert}(S, x)$: add element x to set S .

Maximum (S)

HEAP-MAXIMUM(A)

1 **return** $A[1]$

Costs: $O(1)$

Extract-Max (S)

HEAP-EXTRACT-MAX(A)

```
1  if  $A.heap-size < 1$ 
2      error “heap underflow”
3   $max = A[1]$ 
4   $A[1] = A[A.heap-size]$ 
5   $A.heap-size = A.heap-size - 1$ 
6  MAX-HEAPIFY( $A, 1$ )
7  return  $max$ 
```

Costs: $O(\lg n)$

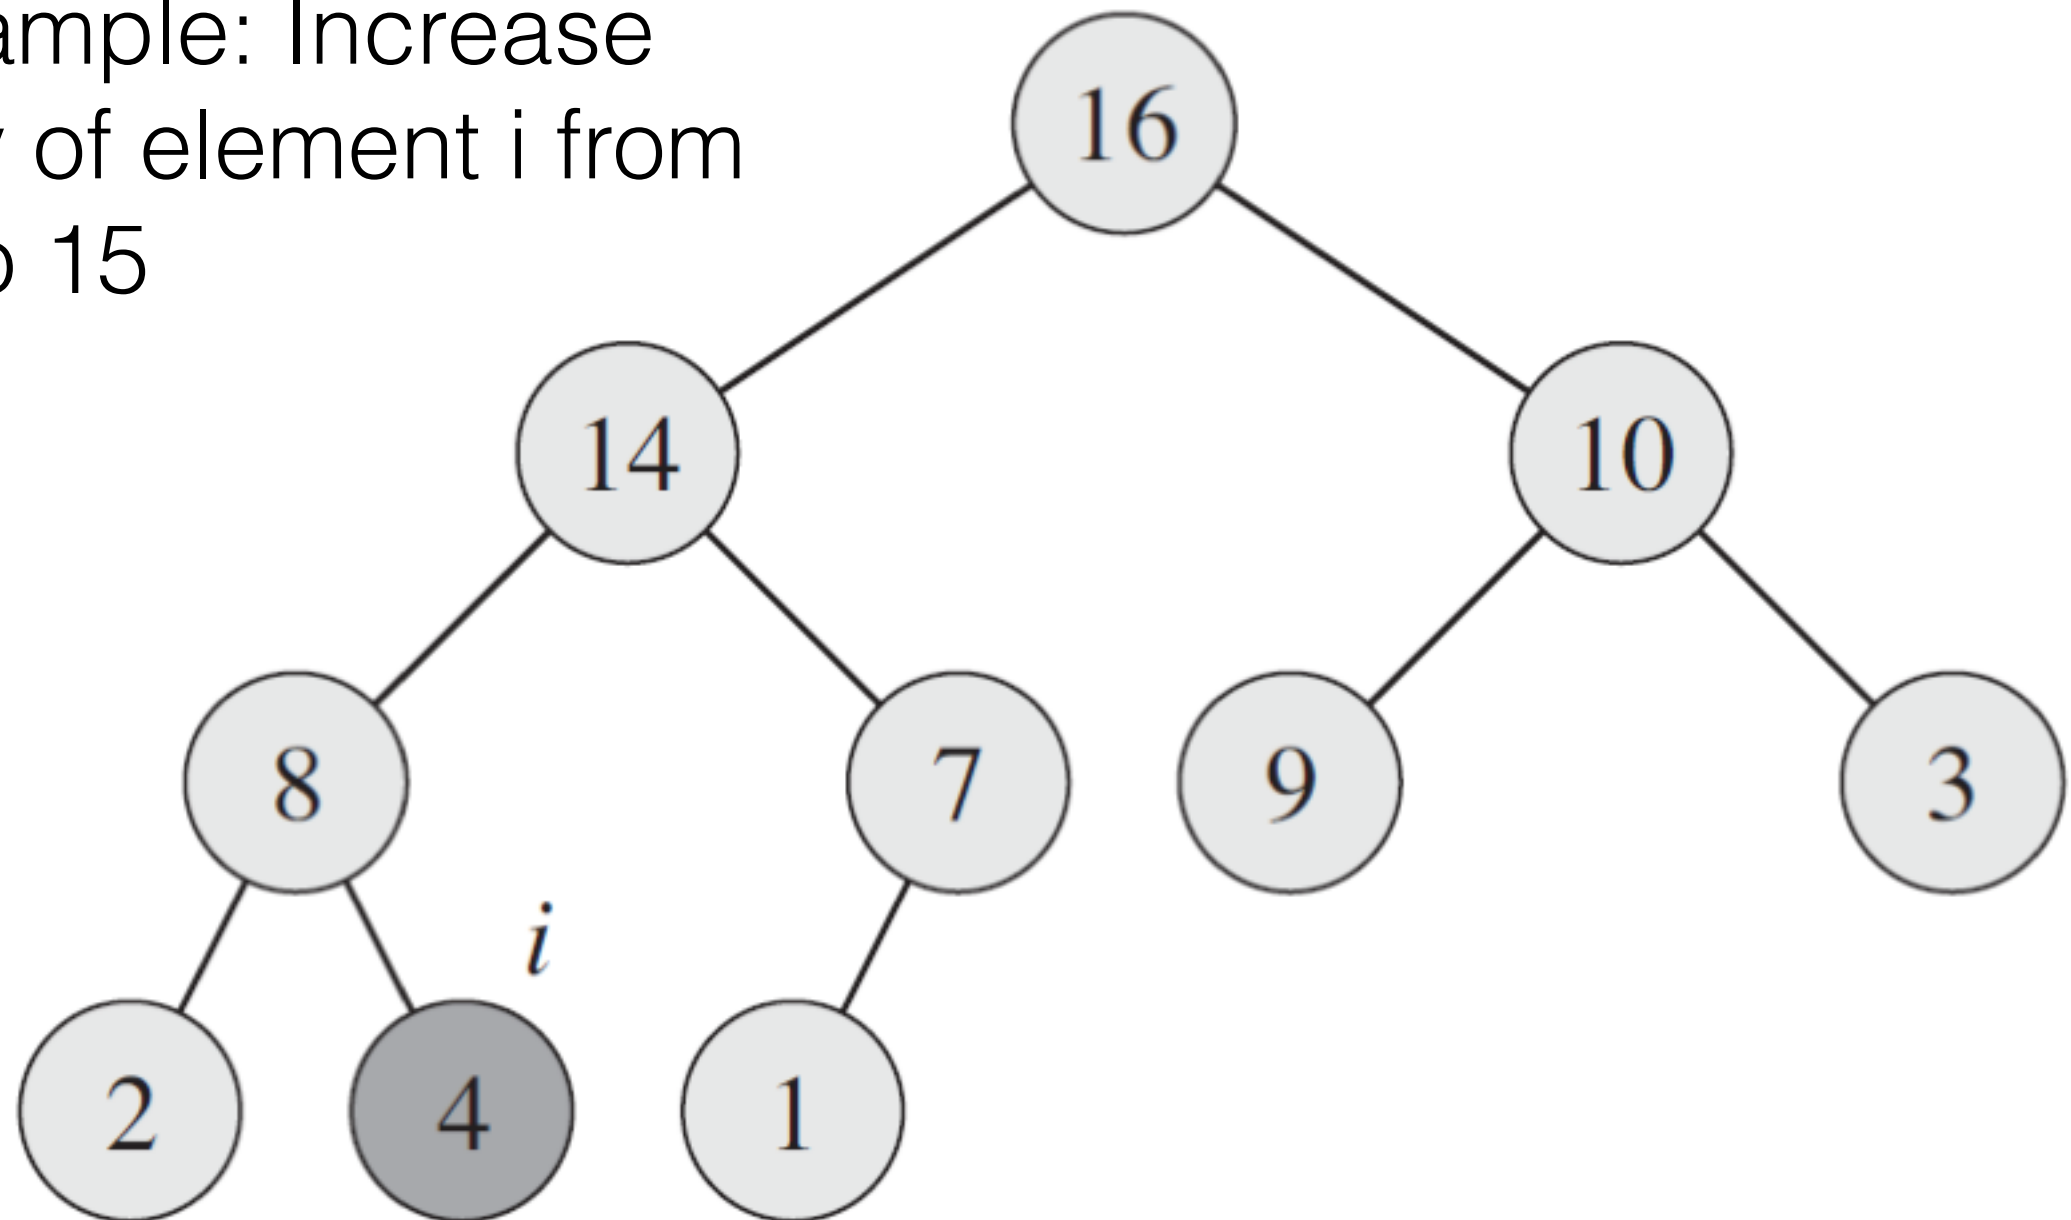
Increase-Key (S, x, k)

HEAP-INCREASE-KEY(A, i, key)


```
1  if  $key < A[i]$ 
2      error “new key is smaller than current key”
3   $A[i] = key$ 
4  while  $i > 1$  and  $A[\text{PARENT}(i)] < A[i]$ 
5      exchange  $A[i]$  with  $A[\text{PARENT}(i)]$ 
6       $i = \text{PARENT}(i)$ 
```

Increase-Key (S, x, k)


Example: Increase
key of element i from
4 to 15




Increase-Key (S, x, k)



Increase-Key (S, x, k)



Increase-Key (S, x, k)



Increase-Key (S, x, k)

HEAP-INCREASE-KEY(A, i, key)

```
1  if  $key < A[i]$ 
2      error “new key is smaller than current key”
3   $A[i] = key$ 
4  while  $i > 1$  and  $A[\text{PARENT}(i)] < A[i]$ 
5      exchange  $A[i]$  with  $A[\text{PARENT}(i)]$ 
6       $i = \text{PARENT}(i)$ 
```

Costs: $O(\lg n)$

Insert (S, x)

MAX-HEAP-INSERT(A, key)

- 1 $A.heap-size = A.heap-size + 1$
- 2 $A[A.heap-size] = -\infty$
- 3 HEAP-INCREASE-KEY($A, A.heap-size, key$)

Costs: $O(\lg n)$